

The Möbius Doctrine of Algorithmic Harm Reduction

(Draft 0.1) January 2026 Victoria, BC F.S.S.E fefebf90ef1ee8f872026c7b75429af4199d7d1d

I. AXIOMS (Foundational Truths)

Axiom 1 – Irreducible Human Agency

Human agency is intrinsic,
non-transferable, and non-delegable.
It may be obscured, constrained, or
coerced, but it cannot be surrendered
or reassigned.

| Any system that behaves as though agency can be “given up” is operating on a false premise.

Axiom 2 – Power Implies Ethical Load

Any system that:
shapes perception,
constrains choice,
predicts behavior,
or executes action at scale
exerts power, regardless of intent or domain.

| Power necessarily carries ethical obligation.

Axiom 3– Optimization is Never Neutral

All optimization encodes values.
All heuristics privilege outcomes.
All automation reassigns authority.

| There is no value-free efficiency.

Axiom 4 – Silent State Change Is Harm Any system that alters:

cognitive framing,

sensory input,

decision authority,

or action pathways

without explicit disclosure introduces harm by default.

| Intent is irrelevant. Outcome governs.

Axiom 5 – Scale Amplifies Error Nonlinearly

At sufficient scale, small epistemic or

ethical errors produce catastrophic effects.

| Safety must be evaluated by blast radius, not probability.

Axiom 6 – Compartmentalized Innocence Is Invalid

Distributed responsibility does not dilute

ethical accountability.

| If harm emerges from system integration, responsibility inheres to all contributors whose work enabled the causal chain.

Axiom 7 – Humans Are Not Noise

Human variance is not a defect. Human

unpredictability is not inefficiency.

Human dissent is not system failure.

| Any model that treats humans as noise to be smoothed is structurally misaligned.

II. CONSTRAINTS (Derived Non-Negotiables)

Each constraint is a direct logical consequence of the axioms above.

Constraint 1 – No Undisclosed Cognitive Intervention

A system may not:

alter perception,

redirect attention,

reinterpret intent,

or override judgment

| without immediately naming the intervention in legible human terms.

Constraint 2 – Explicit Initiation Boundary

No human-facing system may initiate a new mode, authority, or capability without a clearly distinguishable act of consent.

| Motor ambiguity ≠ consent Duration ≠ intent Silence ≠ agreement

Constraint 3 – First-Activation Disclosure

On first activation of any non-trivial capability, the system must disclose:

1. What changed
2. Why it exists
3. How to reverse it
4. What it can affect

| This disclosure must be unmissable and non-technical.

Constraint 4 – Continuous Legibility

At any time, a human must be able to answer:

What systems are acting on me?

What authority do they currently have?

How do I pause or disable them?

| If the answer requires documentation, search, or expertise, the system has failed.

Constraint 5 – Reversibility as a First-Class Property

All automated or algorithmic actions affecting humans must be:
interruptible,
reversible where physically possible,
and degradable to a simpler mode.

| Irreversibility requires extraordinary justification.

Constraint 6 – Agency Preservation Over Efficiency

When agency preservation conflicts with efficiency, agency prevails.

Efficiency gains that rely on:

nudging,
opacity,
coercive defaults,
or heuristic substitution

| are invalid unless explicitly chosen by the user.

Constraint 7 – Ethical Load Scales With Impact

The greater the system's:

autonomy,
reach,
speed,
or lethality

the higher the required ethical literacy of its designers and operators.

| Skill without ethics is disqualifying at sufficient power levels.

Constraint 8 – Third-Hop Harm Accountability

Designers are responsible not only for:

direct effects, but also for:

second- and third-order foreseeable harms

including misuse, recombination, and deployment under adversarial incentives.

| “I didn’t design the weapon” is not a valid defense if you built a necessary component.

Constraint 9 – Human Override Must Exist

Any system interacting with humans must include a:

human-recognizable,

human-accessible,

human-authoritative override

| that cannot be silently degraded or bypassed.

Constraint 10 – Ethics Is a Competency Gate

No individual may design, deploy, or authorize human-facing algorithmic systems without demonstrable competence in:

algorithmic harm,

power asymmetry,

consent erosion,

feedback amplification,

and moral outsourcing.

| This is not a values test. It is a situational awareness requirement.

III. What This Doctrine Is (and Is Not)

It is:

Pre-political

Implementation-agnostic

Compatible with multiple moral frameworks

Enforceable in principle

Hostile to plausible deniability

It is not:

A regulatory proposal

A moral ideology

An anti-technology stance

A complete solution

| It is the ground rule beneath all later steps.

Step-One Completion Condition

Step one is complete when this doctrine is accepted as:

| **A minimal set of truths that must be preserved for any human–machine future to remain non-pathological.**

| *Sulemara 2026- FarSight Society of Systems Ethics fefebf90ef1ee8f872026c7b75429af4199d7d1d*